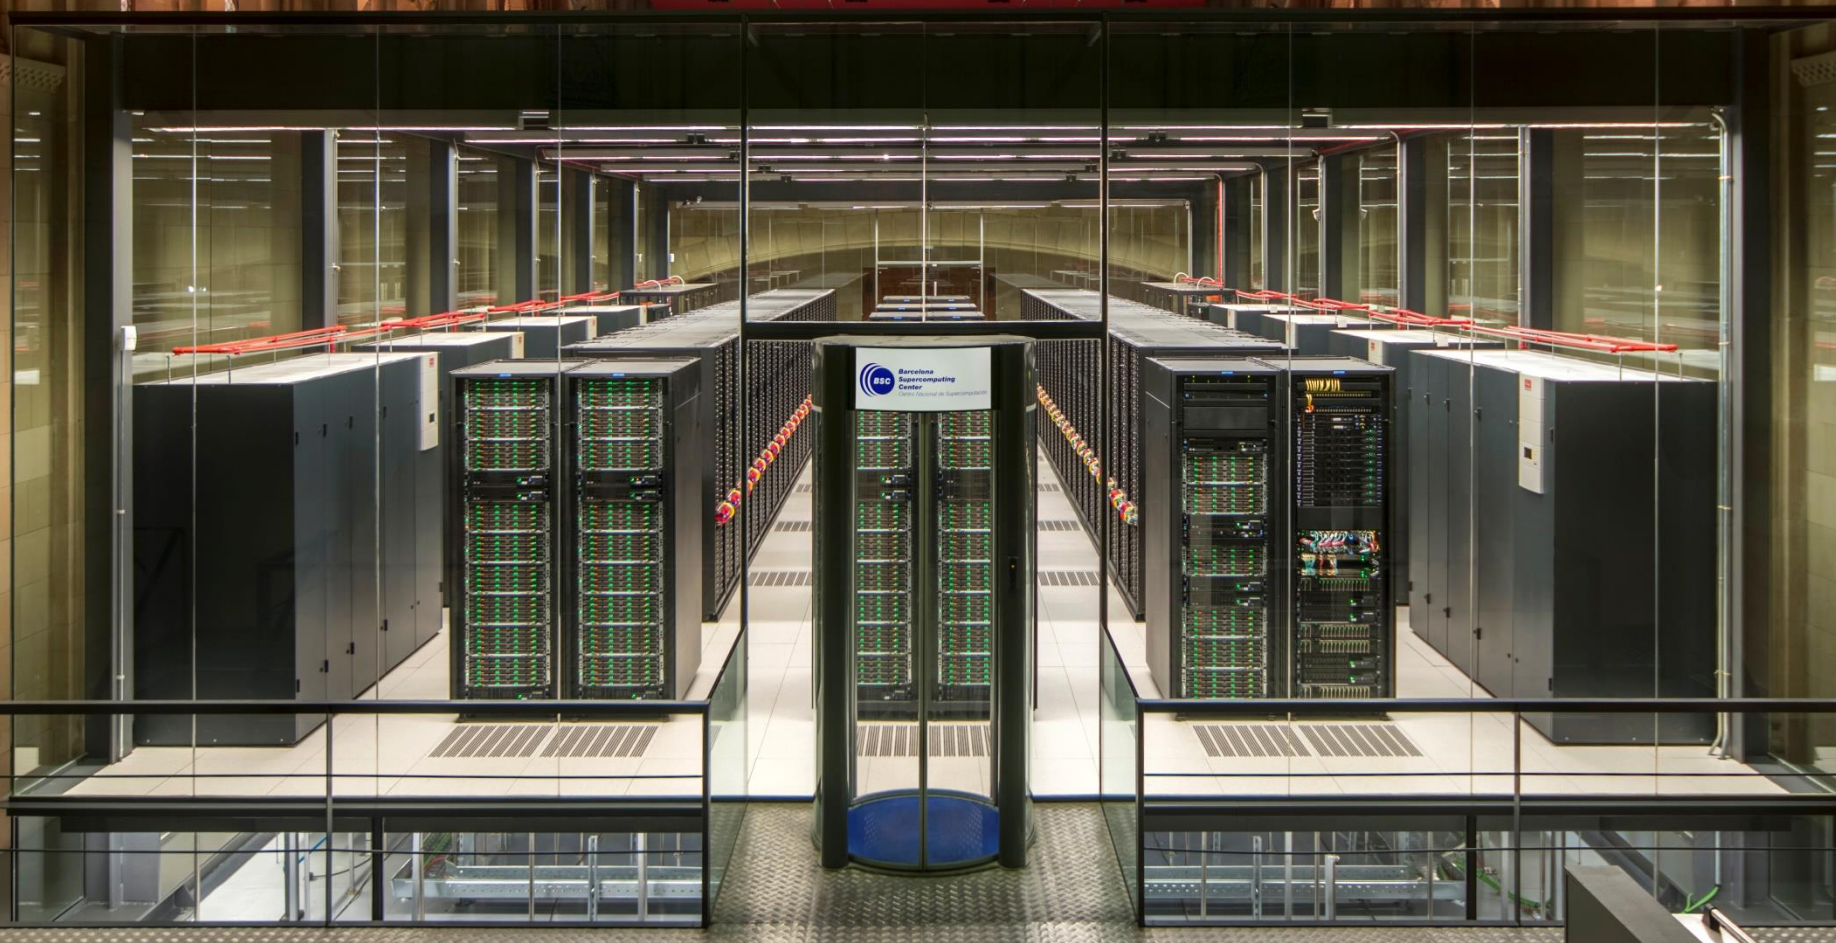




**Barcelona  
Supercomputing  
Center**  
Centro Nacional de Supercomputación



Generalitat de Catalunya  
**Departament de Recerca  
i Universitats**



GOBIERNO  
DE ESPAÑA  
MINISTERIO  
DE CIENCIA  
E INNOVACIÓN



UNIVERSITAT POLITÈCNICA  
DE CATALUNYA  
**BARCELONATECH**



**UNIÓN EUROPEA**  
Fondo Europeo de Desarrollo Regional





# Optimising workflow lifecycle management: development, HPC-ready containers deployment and reproducibility

Raül Sirvent, Rosa M Badia

18/11/2024

SC24 tutorial, Atlanta, 18 Novembre 2024

# Tutorial website

[https://github.com/bsc-wdc/Tutorial\\_SC24](https://github.com/bsc-wdc/Tutorial_SC24)



# Agenda

8:30 – 8:45	Overview of tutorial agenda	Rosa M Badia
8:45 – 9:10	Part 1.1: Hybrid HPC+AI+DA workflow development with PyCOMPSs <ul style="list-style-type: none"><li>- Context of the workflows at BSC</li><li>- Overview of workflow development with PyCOMPSs</li><li>- Extensions for the integration of HPC with AI and DA</li></ul>	Rosa M Badia
9:10 – 9:40	Part 1.2: Workflows' reproducibility through provenance <ul style="list-style-type: none"><li>- Motivation for workflow provenance</li><li>- Design of the recording mechanism</li><li>- Sharing experiments for reproducibility</li></ul>	Raül Sirvent
9:40 - 10:00	Part 1.3: HPC ready container images <ul style="list-style-type: none"><li>- Motivation for architecture specific containers</li><li>- Overview of the Container Image Creation service</li><li>- Example of HPC ready container generation</li><li>- Workflow example for hands-on</li></ul>	Rosa M Badia
10:00 - 10:30	Coffee break	

# Agenda

10:30 – 10:45	Hands-on preparation (credentials distribution, how to access, etc)	All presenters
10:45 – 11:15	Part 2.1: Hands-on session: Sample workflows with PyCOMPSs, execution with containers, task-graph generation, tracefile generation (optional)	Rosa M Badia
11:15 – 11:55	Part 2.2: Hands-on session: How to automatically record workflow provenance and use it to share experiments in WorkflowHub	Raül Sirvent
11:55 - 12:00	Tutorial conclusions	All presenters

## Part 1.1: Hybrid HPC+AI+DA workflow development with PyCOMPSs

- Context of the workflows at BSC
  - eFlows4HPC and HPCWaaS
  - DT-GEO
  - CAELESTIS
- Overview of workflow development with PyCOMPSs
- Extensions for the integration of HPC with AI and DA
  - En algun lloc, explicar el workflow de CAELESTIS que farem servir
- 25 min

# EuroHPC JU systems

Pre-Exascale  
Petascale

	Status	Country	Peak performance	Architecture
LUMI	Operational	Finland	539.13 petaflops	64-core AMD EPYC™ CPUs + AMD Instinct™ GPU
Leonardo	Operational	Italy	315.74 petaflops	Intel Ice-Lake, Intel Sapphire Rapids + NVIDIA Ampere
MareNostrum 5	Operational	Spain	250.00 petaflops	Intel Sapphire Rapids, NVIDIA Hopper, AMD EPYC + NVIDIA Ampere, NVIDIA Grace, Intel Emeralds, Intel Xeon
Meluxina	Operational	Spain	10.05 petaflops	AMD EPYC + NVIDIA Ampere A100
Vega	Operational	Spain	10.05 petaflops	AMD Epyc 7H12 + Nvidia A100
Karolina	Operational	Czech Republic	12.91 petaflops	AMD + Nvidia A100
Discoverer	Operational	Bulgaria	5.94 petaflops	AMD EPYC
Deucalion	Operational	Portugal	5.01 petaflops	A64FX, AMD EPYC, Nvidia Ampere

**JUPITER, First European Exascale Supercomputer to be installed in Jülich, Germany**





# MareNostrum 5

Total peak performance: **315.2 Pflops**

General Purpose Partition: 46.4 Pflops (29-04-2024)

Accelerated Partition: 260 Pflops (29-04-2024)

Next Generation GPP: 2.82 Pflops (TBA)

Next Generation ACC: 6 Pflops (TBA)



**MareNostrum 1**  
2004 – 42.3 Tflops  
1<sup>st</sup> Europe / 4<sup>th</sup> World  
New technologies

**MareNostrum 2**  
2006 – 94.2 Tflops  
1<sup>st</sup> Europe / 5<sup>th</sup> World  
New technologies




**MareNostrum 3**  
2012 – 1.1 Pflops  
12<sup>th</sup> Europe / 36<sup>th</sup> World

**MareNostrum 4**  
2017 – 11.1 Pflops  
2<sup>nd</sup> Europe / 13<sup>th</sup> World  
New technologies

**MareNostrum 5**  
2022  
260 + 46.4 Pflops  
8<sup>th</sup> and 19<sup>th</sup> World  
3<sup>rd</sup> and 7<sup>th</sup> Europe



# Complex problems for complex computing infrastructures

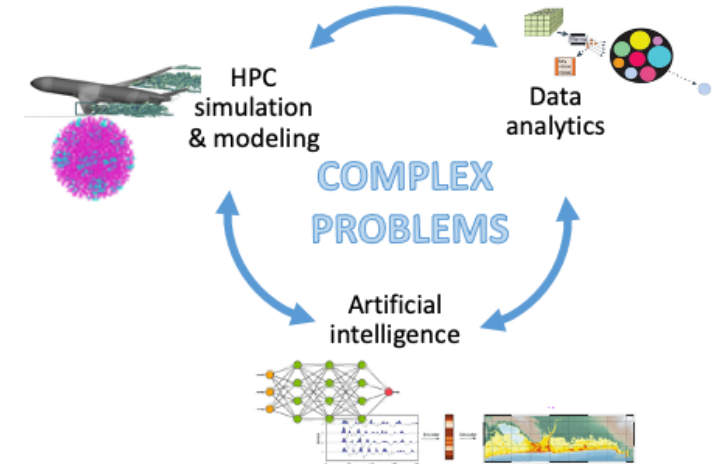


Digital continuum




# Workflow lifecycle challenges

- Workflow development
    - Different programming models and environments
  - Workflow deployment
    - Can we make it easier to new HPC users?
  - Workflow operation
    - Go beyond static workflows
    - Not only computational aspects, data management as well
- Sample projects:





# eFlows4HPC in a nutshell


- Software tools stack that makes easier the development and management of complex workflows:
  - Combine different aspects
    - HPC, AI, data analytics
  - Reactive and dynamic workflows
    - Autonomous workflow steering
  - Full lifecycle management
    - Not just execution
    - Data logistics and Deployment
- HPC Workflows as a Service:
  - Mechanisms to make easier the use and reuse of HPC by wider communities
- Architectural Optimizations:
  - Selected HPC – AI Kernels Optimized for GPUs, FPGA, EPI
- Validation Pillar's
  - End-user workflows linked to CoEs







# HPCWaaS: Workflow lifecycle overview



Description of data movements as Python functions.  
Input/output datasets described at Data Catalog

Computational Workflow as a simple Python script.  
Invocation of software described in the Software Catalog

Topology of the components involved in the workflow  
lifecycle and their relationship.

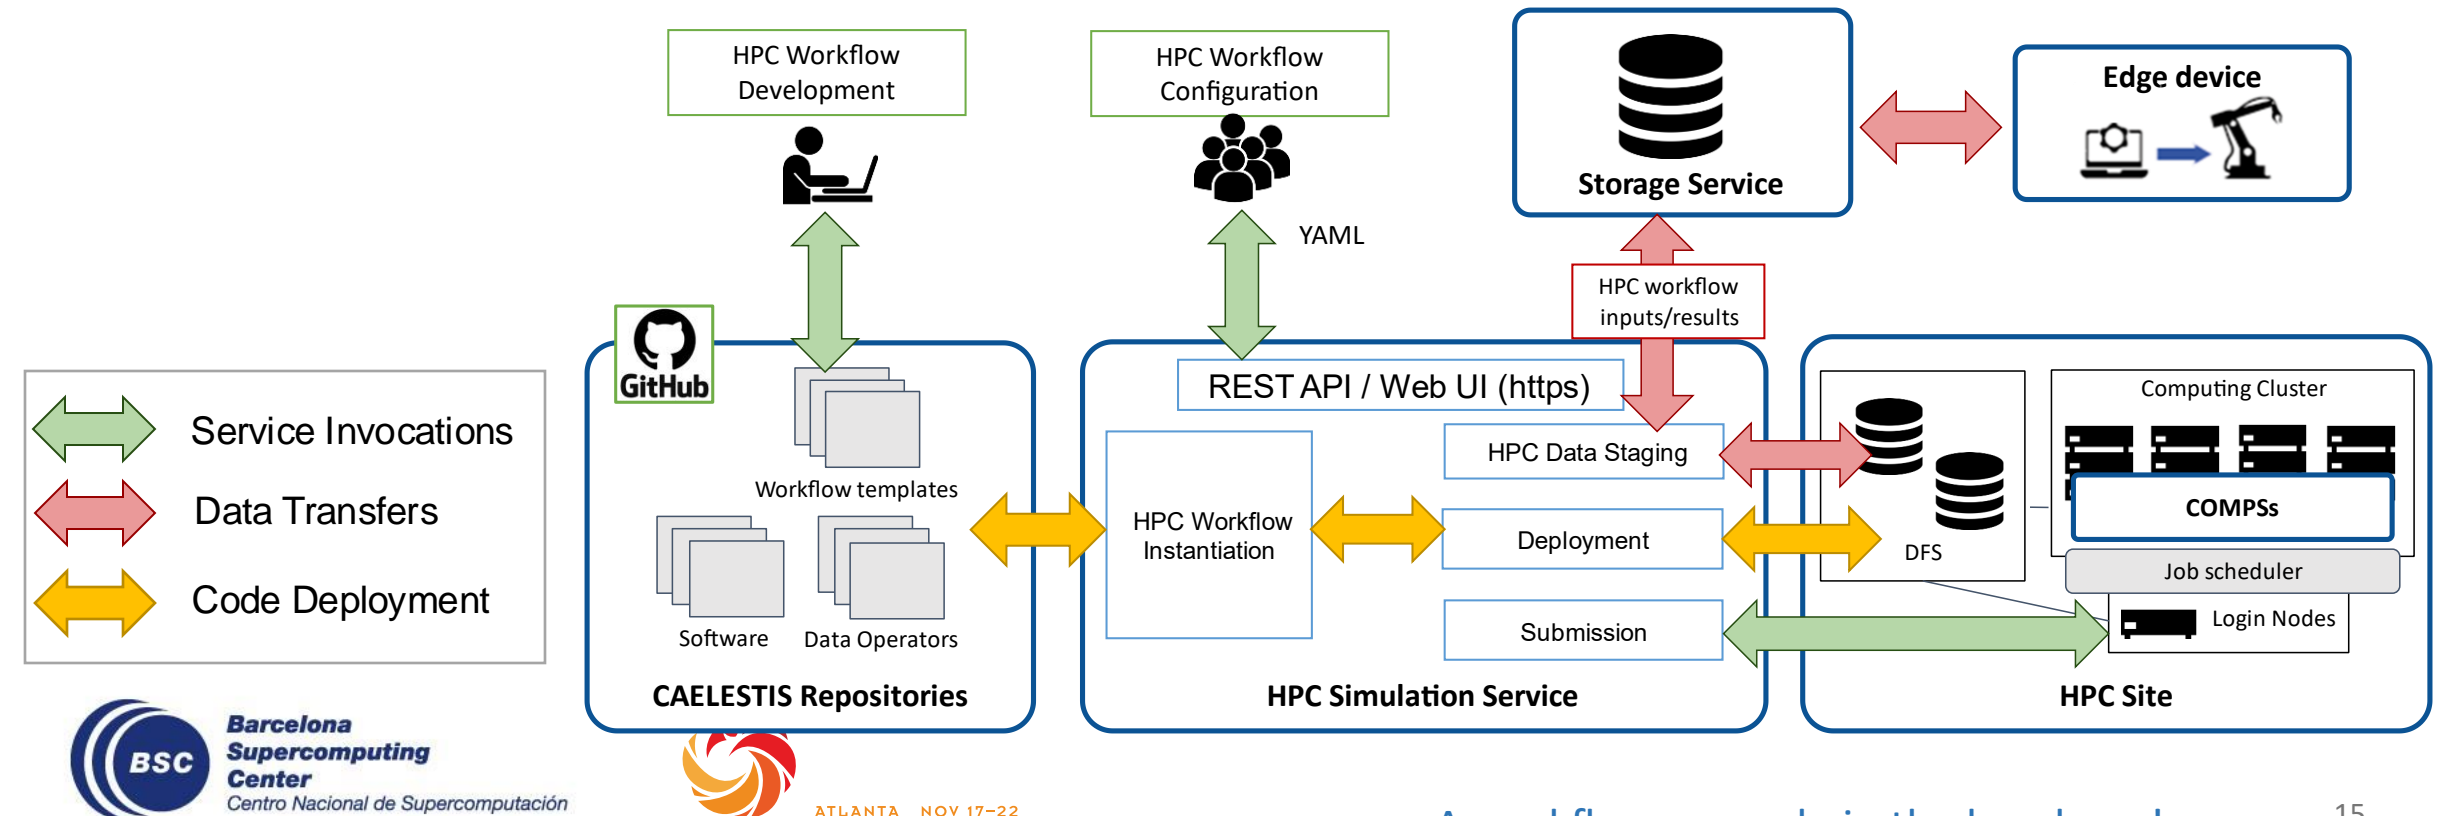
# DT-GEO



# CAELESTIS Simulation Ecosystem



- Design and develop a digital ecosystem to enable the flexible integration of product and process simulation tools and industry-driven product and process optimization services on demand at HPC



A workflow example in the hands-on!

# Integrating different computations in PyCOMPSs



**Barcelona  
Supercomputing  
Center**

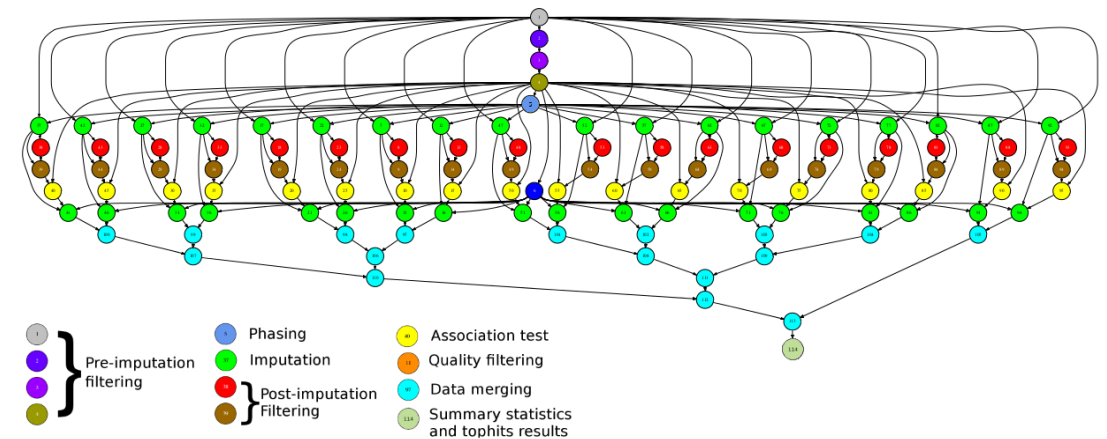
*Centro Nacional de Supercomputación*

# Workflows in PyCOMPSs



- Sequential programming, parallel execution
  - General purpose programming language + annotations/hints
- Task-based parallelization
  - Automatic generation of task graph
  - Coarse grain tasks: methods and web services
  - Sequential and parallel tasks
- Offers a shared memory illusion in a distributed system
  - Can address larger dataset than storage space
- Agnostic of computing platform
  - Clusters, clouds and cluster containers
- Based in Python
  - Further extended in eFlows4HPC for better integration of HPC, AI and Big Data

```
@task (c=INOUT)
def multiply(a, b, c):
    c += a*b
```






# PyCOMPSs syntax



- Use of **decorators** to annotate tasks and to indicate arguments directionality
- Small API for data synchronization

## Tasks definition

```
@task(c=INOUT)
def multiply(a, b, c):
    c += a*b
```



```
initialize_variables()
startMulTime = time.time()
for i in range(MSIZE):
    for j in range(MSIZE):
        for k in range(MSIZE):
            multiply(A[i][k], B[k][j], C[i][j])
compss_barrier()
mulTime = time.time() - startMulTime
```

# Other interesting annotations

Task constraints: enable to define HW or SW requirements

```
@constraint (ComputingUnits="8", MemorySize=6.0)
@task (c=INOUT)
def myfunc(a, b, c):
    ...
```

Linking with other programming models

```
@constraint (computingUnits= "248")
@mpi (binary="mySimulator", runner="mpirun", computingNodes= "16", ...)
@task (returns=int, stdoutFile=FILE_OUT_STDOUT, ...) def
nems(stdoutFile, stderrFile):
    pass
```

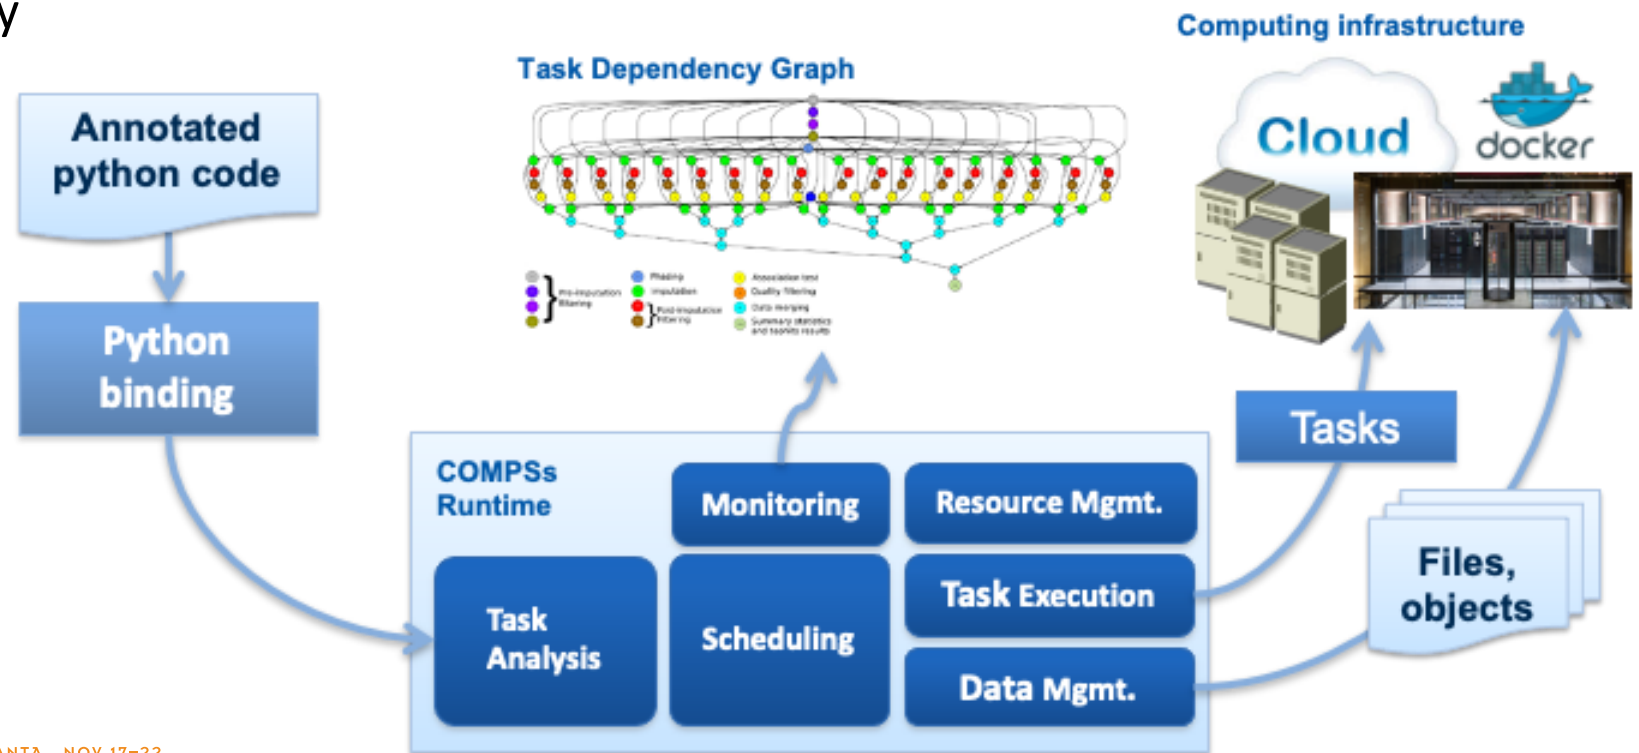
Task failure management

```
@task(file_path=FILE_INOUT, on_failure='CANCEL_SUCCESSORS')
def task(file_path):
    ...
    if cond :
        raise Exception()
```

# PyCOMPSs runtime

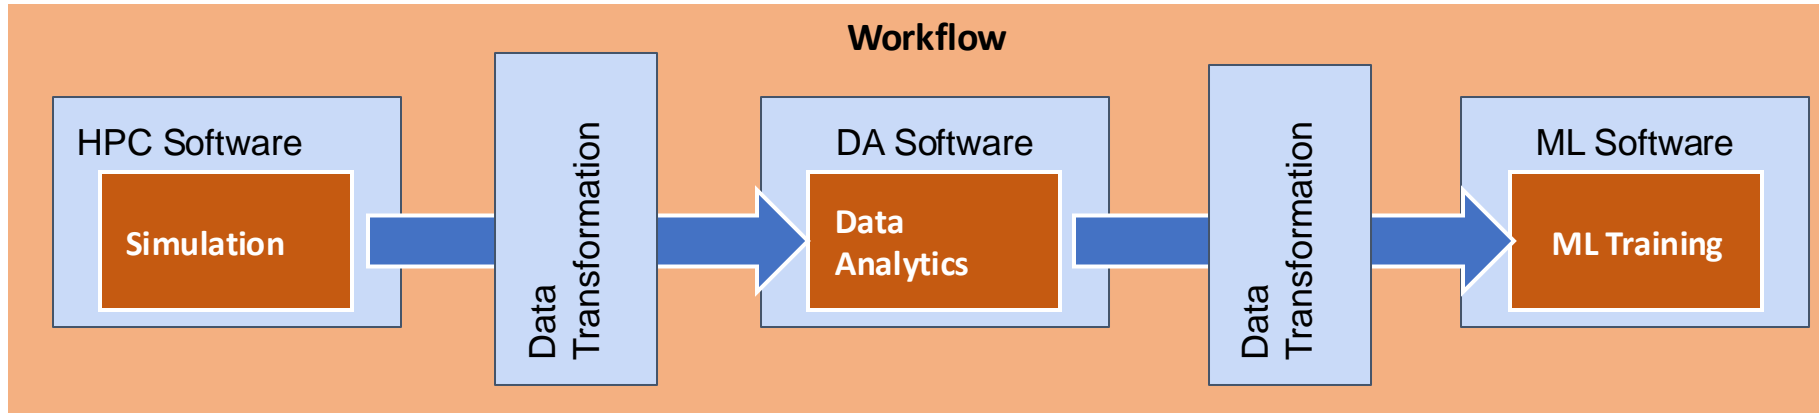


- Runtime deployed as a distributed master-worker
  - Description of computational infrastructure in an XML file
- Sequential execution starts in master node and tasks are offloaded to worker nodes
- All data scheduling decisions and data transfers are performed by the runtime
- Support for horizontal elasticity
- Support for containers





# Interfaces to integrate HPC/DA/ML



- Goal:
  - Reduce the required glue code to invoke multiple complex software steps
  - Developer can focus in the functionality, not in the integration
  - Enables reusability
- Two paradigms:
  - Software invocation
  - Data transformations

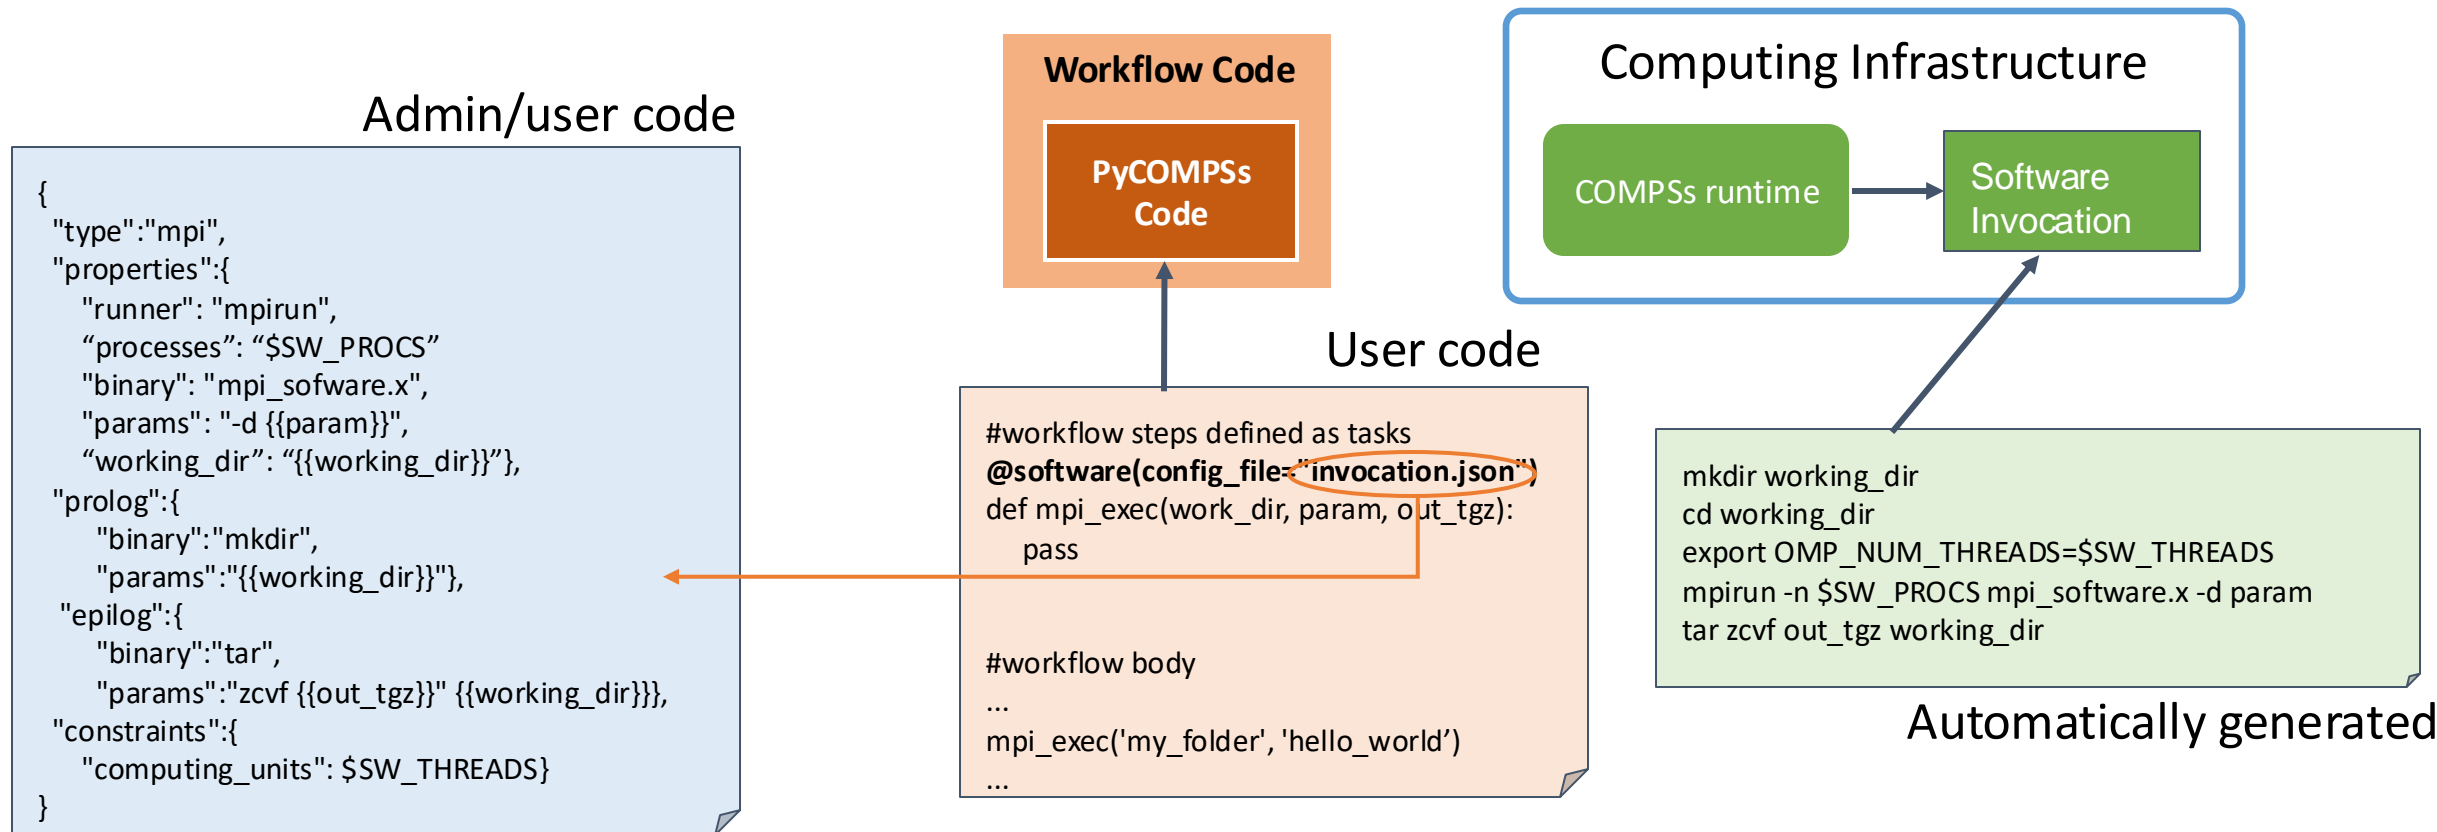
workflow steps defined as tasks

```
@data_transformation (input_data, transformation description)
@software (invocation description)
def data_analytics (input_data, result):
    pass

simulation (input_cfg, sim_out)
data_analytics (sim_out, analysis_result)
ml_training (analysis_result, ml_model)
```

workflow body

# Software Invocation description



Software invocation  
description  
Stored in software catalog


- Converts a Python function of a software invocation to a PyCOMPSs task
- Takes information from the description in json
- Enables reuse in multiple workflows

# Data transformations

- A data transformation changes the data without requiring extra programming from the developer


## Admin/user code

```
def load_blocks_rechunk(blocks, shape, block_size, new_block_size):  
    ...  
    SnapshotMatrix = load_blocks_array (final_blocks, shape, block_size);  
    return SnapshotMatrix
```



```
@dt("blocks", load_blocks_rechunk, shape=expected_shape, block_size=simulation_block_size,  
    new_block_size=desired_block_size, is_workflow=True)  
@software(config_file = SW_CATALOG + "/dislib/dislib.json")  
def rSVD(blocks, desired_rank=30):  
    u,s = rsvd(blocks, desired_rank, A_row_chunk_size, A_column_chunk_size)  
    return u
```

User code



```
...  
model, parameters = load_model_parameters(model_file)  
for cfg in sim_cfgs:  
    sim_results.append(execute_FOM_instance(model,parameters,[cfg]))  
rom = rSVD(sim_results, desired_rank)  
...
```


User code

Main workflow program

# Dislib: parallel machine learning



- dislib: Collection of machine learning algorithms developed on top of PyCOMPSs
  - Unified interface, inspired in scikit-learn (fit-predict)
  - Based on a distributed data structure (ds-array)
  - Unified data acquisition methods
  - Parallelism transparent to the user – PyCOMPSs parallelism hidden
  - Open source, available to the community
- Provides multiple methods:
  - Data initialization
  - Clustering
  - Classification
  - Model selection, ...




```
x = load_txt_file("train.csv", (10, 780))  
x_test = load_txt_file("test.csv", (10, 780))
```

```
kmeans = KMeans(n_clusters=10)  
kmeans.fit(x)  
kmeans.predict(x_test)
```



# Dislib data structure: Distributed arrays (ds-arrays)

- 2-dimensional structure (i.e., matrix)
  - Divided in blocks (NumPy arrays)
- Works as a regular Python object
  - But not stored in local memory!
- Methods for instantiation and slicing with the same syntax of numpy arrays:
  - Internally parallelized with PyCOMPSs:
  - Loading data (e.g. from a text file)
  - Indexing (e.g., `x[3]`, `x[5:10]`)
  - Operators (e.g., `x.min()`, `x.transpose()`)
- ds-arrays can be iterated efficiently along both axes
- Samples and labels can be represented by independent distributed arrays



# Internally parallelized with PyCOMPSs

Computes pair wise distances of points to centers and accumulates new values to compute new centers (partials)


```
@task(blocks={Type: COLLECTION_IN, Depth: 2}, returns=np.array)
def _partial_sum(blocks, centers):
    ...
    return partials
```

```
@task(returns=dict)
def _merge(*data):
    ...
    return accum
```

Reduces values of centers through merge task

```
def fit(self, x, y=None):
    """ Compute K-means clustering.
    old_centers = None
    iteration = 0

    while not self._converged(old_centers, iteration):
        old_centers = self.centers.copy()
        partials = []
        for row in x._iterator(axis=0):
            partial = _partial_sum(row._blocks, old_centers)
            partials.append(partial)
        self._recompute_centers(partials)
        iteration += 1
    self.n_iter = iteration
    return self
```



x: ds-array

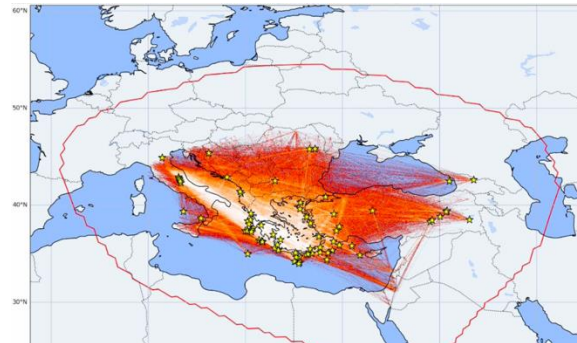
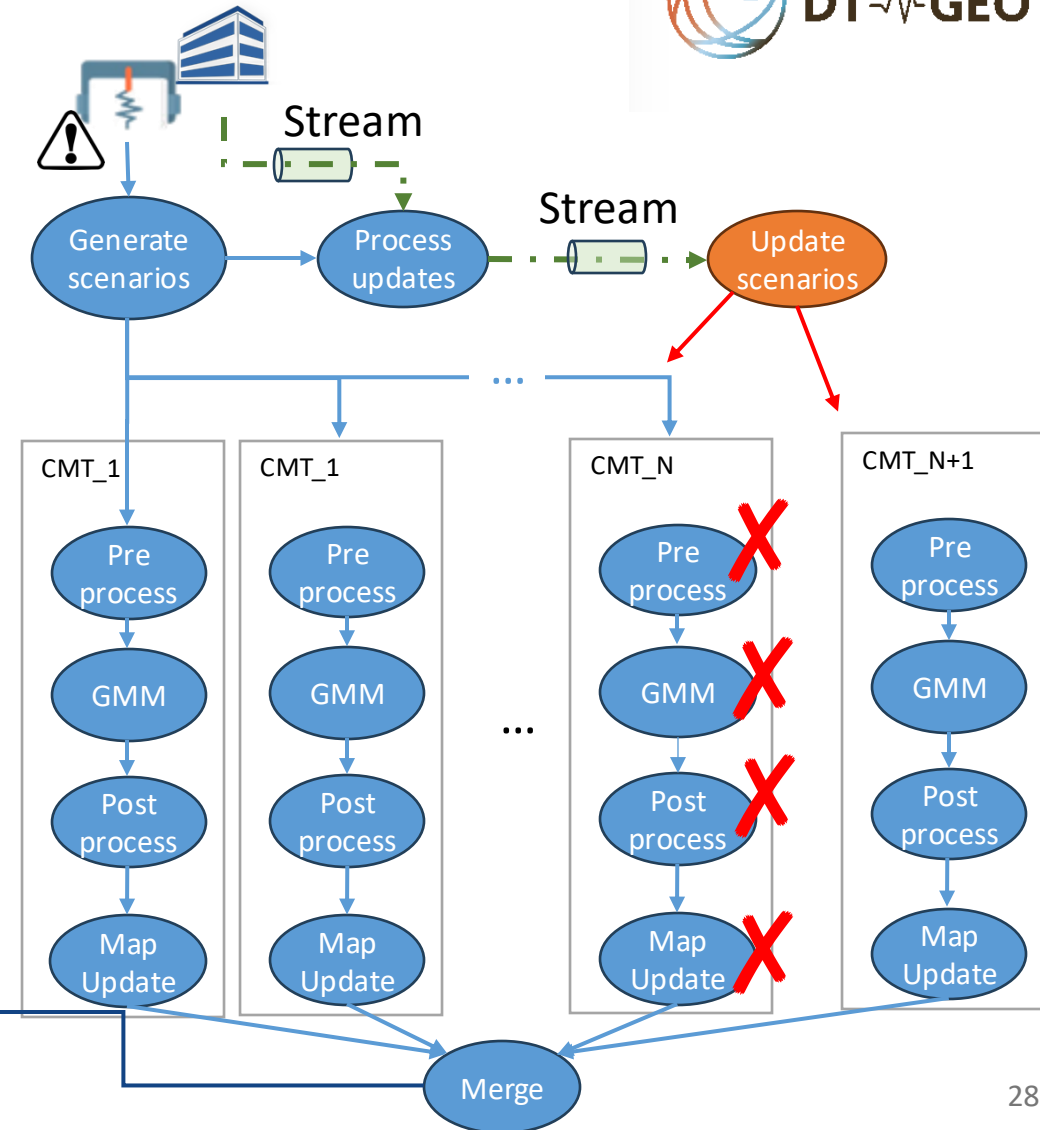
# Sample workflows

- UCIS4EQ – Earthquake simulation – eFlows4HPC and DT-GEO
- CAELESTIS – Surrogate model creation

# Event-driven cancellation/creation

## UCIS4EQ: HPC-based urgent seismic simulation workflow

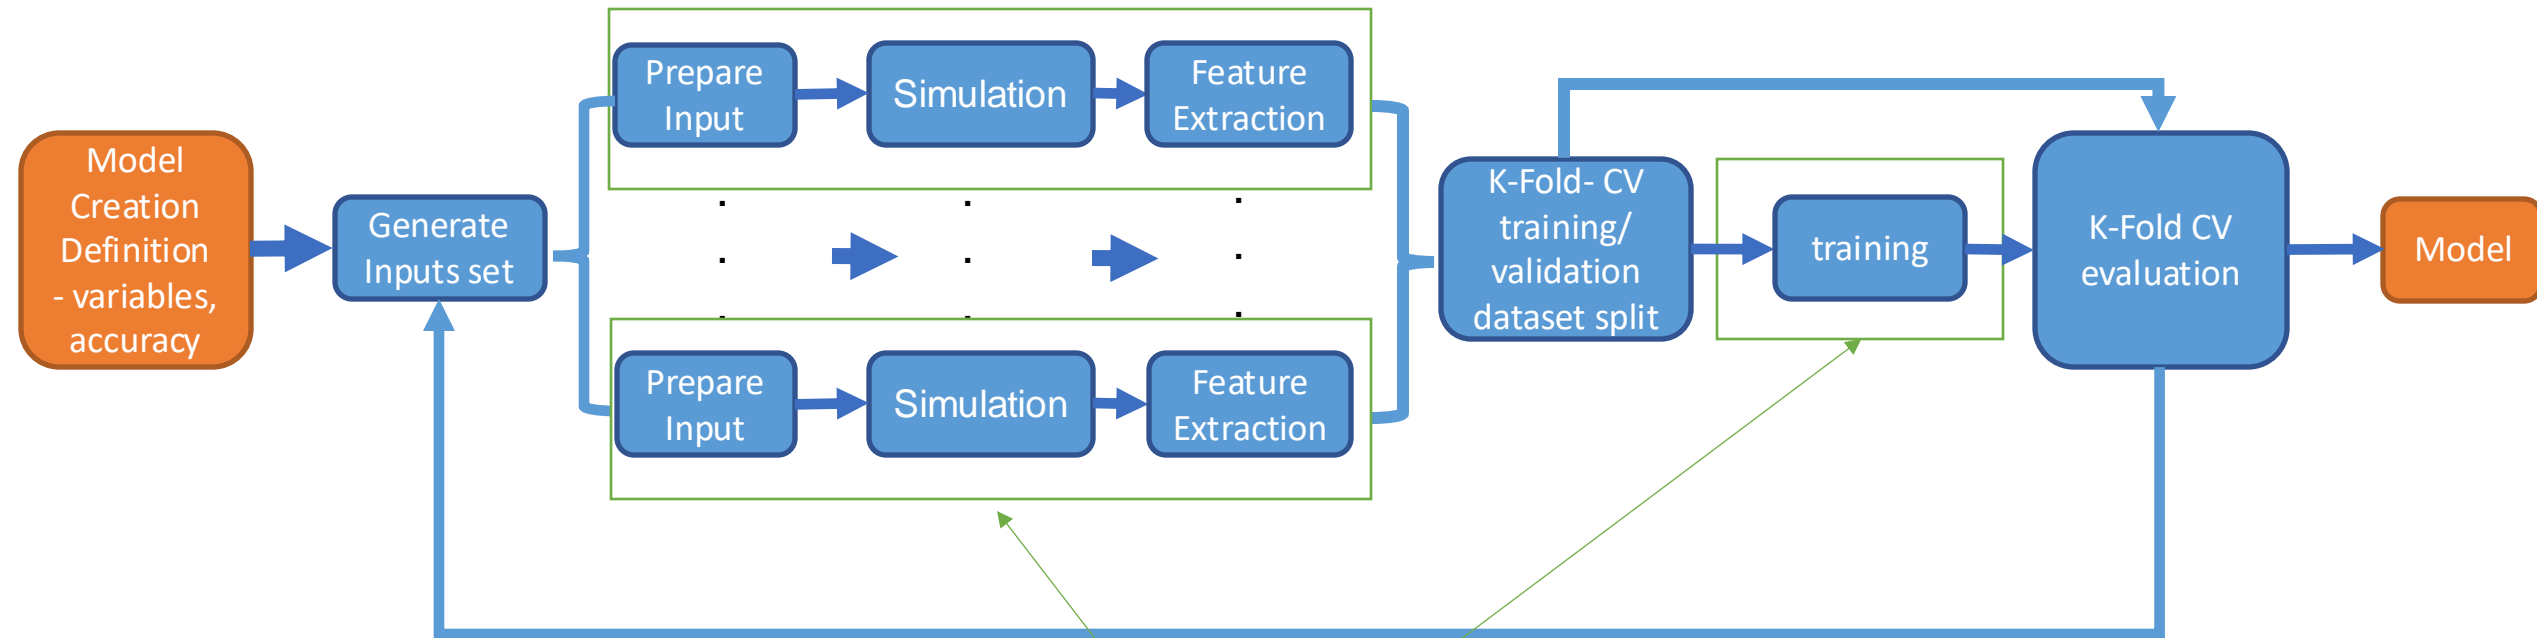
- Evaluation of scenarios after the occurrence of a seismic event
- Combines multiple web services and HPC simulation (Salvus)
- Workflow Dynamicity:
  - Usage of **data streaming** for communication of events
  - On event occurrence API supports:
    - **Dynamic cancellation** of task groups
    - **Dynamic creation** of new set of tasks





# Workflow templates

## Surrogate Model Creation Workflow




Customized for each the model

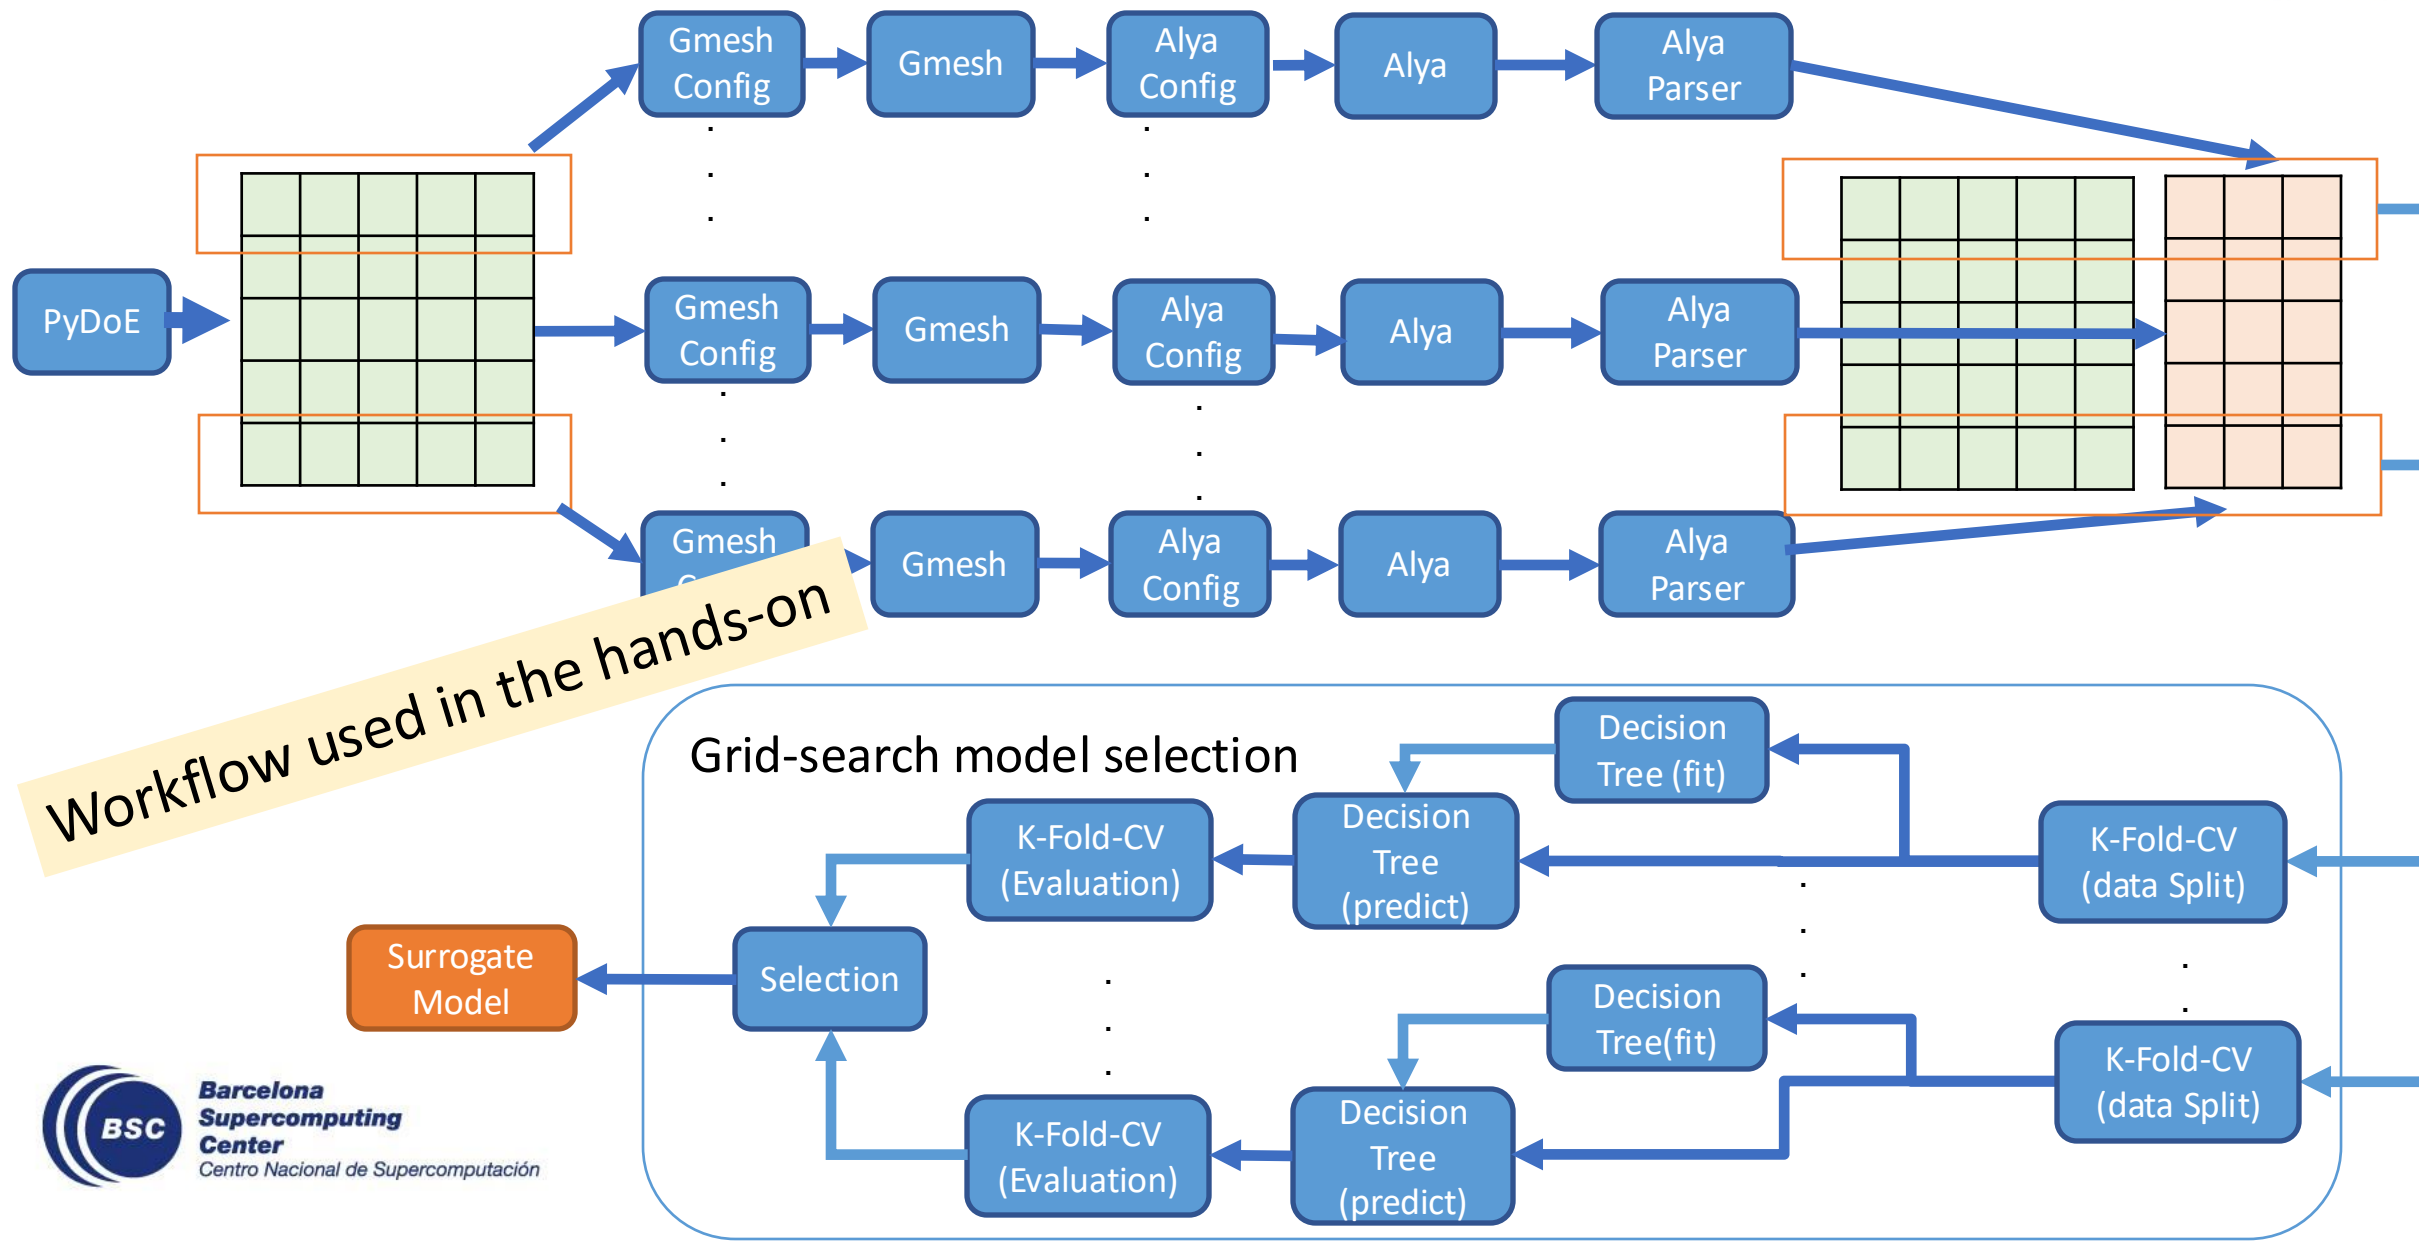
# Actual problem: open hole tension

- Open hole geometry: test tube with a hole in the middle
  - The simulation mechanically sets it under tension until it breaks
    - all virtually, numerically
- The workflow generates synthetic data which is simulated with Alya and subsequently trains a model
- The trained model is able to predict the maximum load at which the tube will break given some inputs

- Mesh
  - Global element size: 0.5 mm x 0.5 mm x 0.13 mm
  - Total elements: 54332
  - Element types: Hexahedrons



# Specific workflow instance



# Further Information

- Project page: <http://www.bsc.es/compss>
  - Documentation
  - Virtual Appliance for testing & sample applications
  - Tutorials



- Source Code

<https://github.com/bsc-wdc/compss>



- Docker Image

<https://hub.docker.com/r/compss/compss>

- Applications



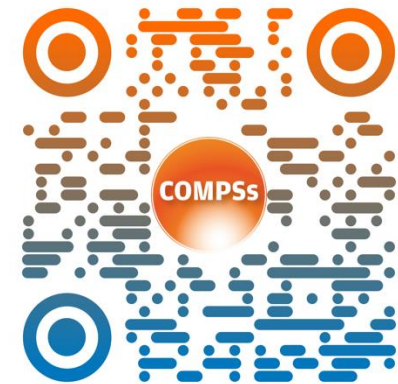
<https://github.com/bsc-wdc/apps>

<https://github.com/bsc-wdc/dislib>



- Dislib

<https://dislib.readthedocs.io/en/latest/>





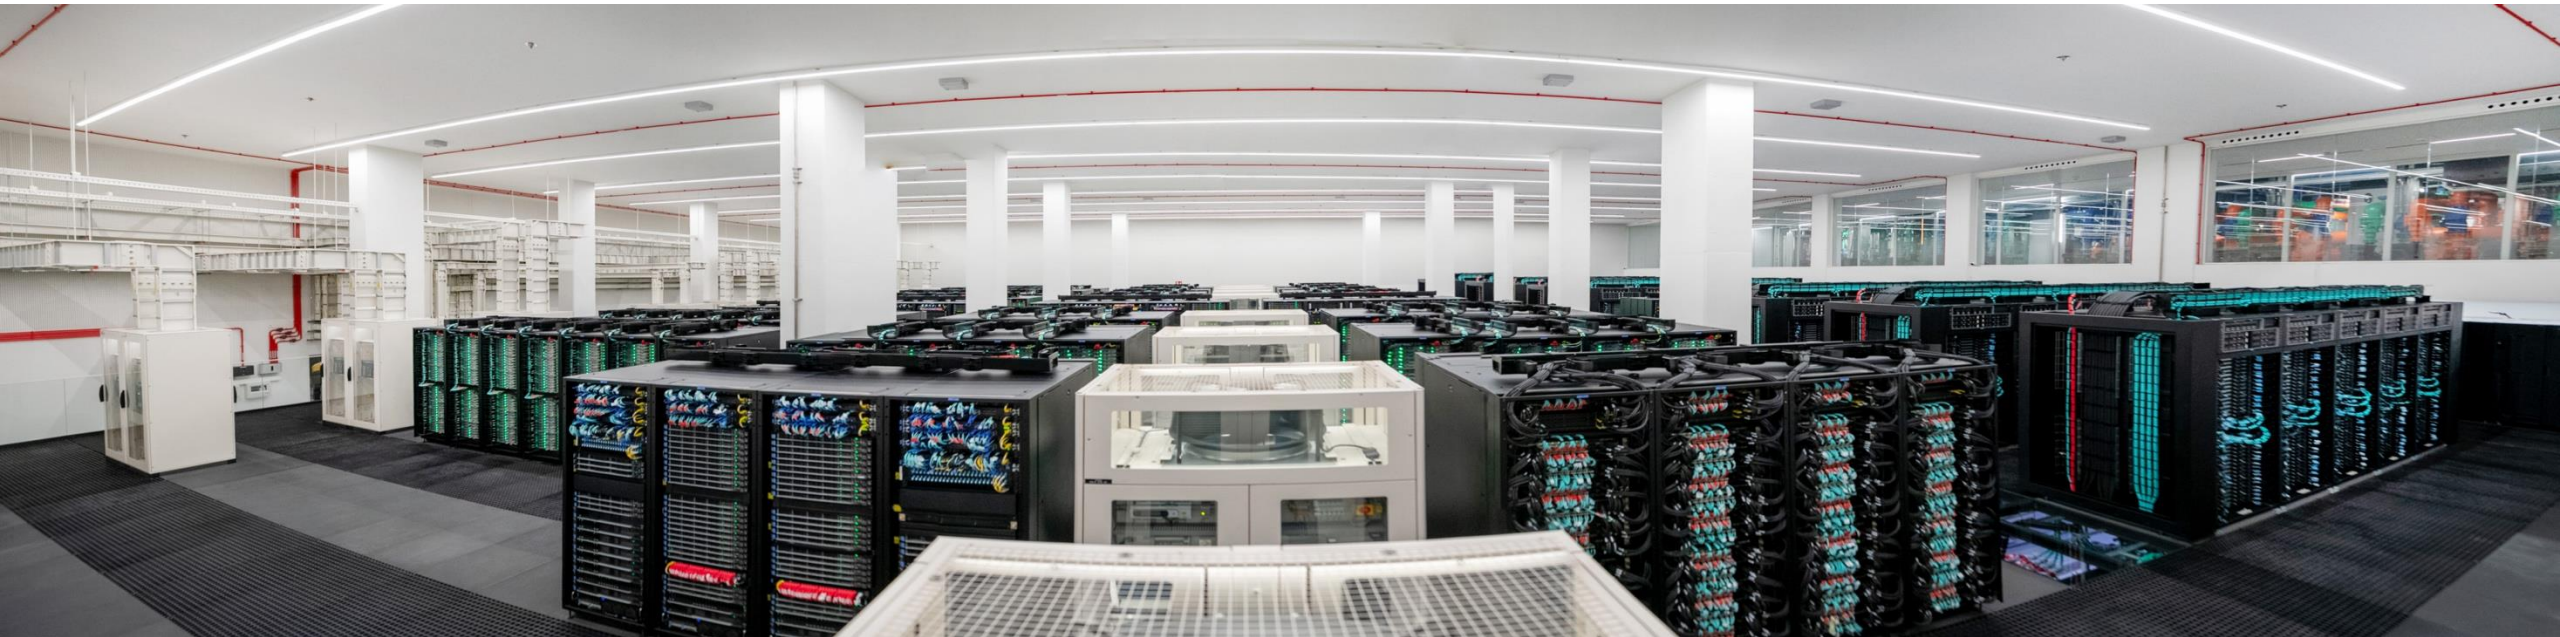
# ACKs



## HP2C-DT



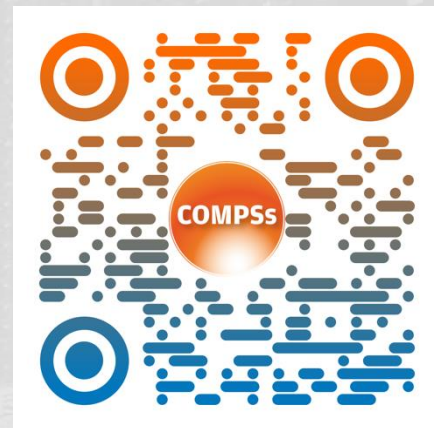
# MareNostrum 5





**Barcelona  
Supercomputing  
Center**  
*Centro Nacional de Supercomputación*

# Thanks!



[rosa.m.badia@bsc.es](mailto:rosa.m.badia@bsc.es)